

✓ Lab 4: Data Visualization and EDA

CPE232 Data Models

1. Load all Superstore datasets.

Note: The same datasets used in Lab 3

```
#!/pip install missingno
#!/pip install plotly
#!/pip install -upgrade nbformat
```

```
# Write your code here
import pandas as pd
import numpy as np
import missingno as msno
```

```
df_superstore_order = pd.read_csv('superstore_order.csv')
df_superstore_people = pd.read_csv('superstore_people.csv')
df_superstore_reutrnr = pd.read_csv('superstore_return.csv')
```

2. Determine shape of each dataset (print out the results as well).

```
df_superstore_order.shape, df_superstore_people.shape, df_superstore_reutrnr.shape
```

```
→ ((8880, 21), (4, 2), (296, 2))
```

3. Show information of the dataset.

```
# Write your code here
print(df_superstore_order.info())
```

```
→ <class 'pandas.core.frame.DataFrame'>
RangeIndex: 8880 entries, 0 to 8879
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Row ID                 8880 non-null  int64
1   Order ID              8880 non-null  object
2   Order Date            8880 non-null  object
3   Ship Date             8880 non-null  object
4   Ship Mode             8880 non-null  object
5   Customer ID           8880 non-null  object
6   Customer Name         8880 non-null  object
7   Segment              8880 non-null  object
8   Country               8880 non-null  object
9   City                 8880 non-null  object
10  State                8880 non-null  object
11  Postal Code          8880 non-null  int64
12  Region              8880 non-null  object
13  Product ID           8880 non-null  object
14  Category             8880 non-null  object
15  Sub-Category         8880 non-null  object
16  Product Name         8880 non-null  object
17  Sales                8880 non-null  float64
18  Quantity             8880 non-null  int64
19  Discount             8880 non-null  float64
20  Profit               8880 non-null  float64
dtypes: float64(3), int64(3), object(15)
memory usage: 1.4+ MB
None
```

```
print(df_superstore_people.info())
```

```
→ <class 'pandas.core.frame.DataFrame'>
RangeIndex: 4 entries, 0 to 3
Data columns (total 2 columns):
#   Column  Non-Null Count  Dtype
---  -
0   Person  4 non-null       object
1   Region  4 non-null       object
```

```
dtypes: object(2)
memory usage: 196.0+ bytes
None
```

```
print(df_superstore_reutrnrn.info())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 296 entries, 0 to 295
Data columns (total 2 columns):
#   Column      Non-Null Count  Dtype
---  ---
0    Returned    296 non-null    object
1    Order ID    296 non-null    object
dtypes: object(2)
memory usage: 4.8+ KB
None
```

4. Are there any missing values? If so, in which column?

Ans: No, there have no any missing values in the dataset.

5.

- 5.1 List unique segments
- 5.2 List unique segments and their corresponding count
- 5.3 Create a pie chart to demonstrate unique segments and their count
- 5.4 Briefly describe what could be interpreted from this pie chart

Note: please create additional cells to answer 5.2 - 5.3

```
# Write your code here (5.1)
# 5.
# - 5.1 List unique segments
# - 5.2 List unique segments and their corresponding count
# - 5.3 Create a pie chart to demonstrate unique segments and their count
# - 5.4 Briefly describe what could be interpreted from this pie chart

# *Note: please create additional cells to answer 5.2 - 5.3*
df_superstore_order['Segment'].unique()
```

```
array(['Consumer', 'Corporate', 'Home Office'], dtype=object)
```

```
# 5.2 List unique segments and their corresponding count
df_superstore_order['Segment'].value_counts()
```

```
count
Segment
Consumer    4613
Corporate    2673
Home Office  1594
```

```
dtype: int64
```


```
# (5.3) Create a pie chart to demonstrate unique segments and their count
import plotly.express as px
```

```
labels = df_superstore_order['Segment'].unique()
data = df_superstore_order['Segment'].value_counts()[labels]
```

```
fig = px.pie(values=data, names=labels, title='Unique Segments and their count')
fig.show()
```



Unique Segments and their count



Answer for the question 5.4

Ans: The pie chart shows that the Consumer segment has the highest count, followed by Corporate and Home Office segments.

- 6.
- 6.1 List unique states
 - 6.2 List top-10 unique states and their corresponding count
 - 6.3 Create a bar chart (vertical) to demonstrate the count of top-10 unique states
 - 6.4 Based on 6.2, also include the total sales of these states (show your result as a dataframe)
 - 6.5 Using the result from 6.4, if you were the owner of this superstore, what information could be interpreted from this result?

Note: please create additional cells to answer 6.2 - 6.4

Write your code here (6.1)

```
df_superstore_order['State'].unique()
```

```
array(['Kentucky', 'California', 'Florida', 'North Carolina',
      'Washington', 'Texas', 'Wisconsin', 'Utah', 'Nebraska',
      'Pennsylvania', 'Illinois', 'Minnesota', 'Michigan', 'Delaware',
      'Indiana', 'New York', 'Arizona', 'Virginia', 'Tennessee',
      'Alabama', 'South Carolina', 'Oregon', 'Colorado', 'Iowa', 'Ohio',
      'Missouri', 'Oklahoma', 'New Mexico', 'Louisiana', 'Connecticut',
      'New Jersey', 'Massachusetts', 'Georgia', 'Nevada', 'Rhode Island',
      'Mississippi', 'Arkansas', 'Montana', 'New Hampshire', 'Maryland',
      'District of Columbia', 'Kansas', 'Vermont', 'Maine',
      'South Dakota', 'Idaho', 'North Dakota', 'Wyoming',
      'West Virginia'], dtype=object)
```

Write your code here (6.2)

```
t10 = df_superstore_order['State'].value_counts().sort_values(ascending=False)[:10].reset_index()
```

```
print(t10)
```


```
State  count
0   California  1754
1    New York  1001
2     Texas    860
3  Pennsylvania  531
4   Washington  452
5    Illinois  427
6     Ohio    396
7    Florida   339
8    Michigan  230
9  North Carolina  229
```

```
# Write your code here (6.3)
import plotly.express as px
```

```
fig = px.bar(t10, x='State', y='count', title='Top 10 States by Order')
fig.show()
```



Top 10 States by Order



```
# Write your code here (6.4)
# include the total sales of these states (show your result as a dataframe)
t10_state = t10['State'].tolist()
t10_state_sales = df_superstore_order[df_superstore_order['State'].isin(t10_state)].groupby('State')['Sales'].sum().reset_index()
t10_state_sales
```



	State	Sales	
0	California	399195.4555	
1	Florida	84083.0880	
2	Illinois	71456.1780	
3	Michigan	62147.6960	
4	New York	274866.8190	
5	North Carolina	49962.1580	
6	Ohio	67924.2140	
7	Pennsylvania	103852.5210	
8	Texas	147855.0282	
9	Washington	124497.7780	

Next steps: [Generate code with t10_state_sales](#) [View recommended plots](#) [New interactive sheet](#)

Answer for the question 6.5

Ans: from the top 10 state total sales owners can see the highest sales state and lowest sales state and plan some strategy to increase the sales in the lowest sales state or can increase the sales in the highest sales state to increase the profit.

7.

- 7.1 List unique categories
- 7.2 Create a bar chart (horizontal) to demonstrate the proportion of these categories
- 7.3 Compute the ratio of these categories in percentage and print the results

Note: please create additional cells to answer 7.2 - 7.3

Write your code here (7.1)

```
df_superstore_order['Category'].unique()
```

```
array(['Furniture', 'Office Supplies', 'Technology'], dtype=object)
```

Write your code here (7.2)


```
category_count = df_superstore_order['Category'].value_counts().reset_index()
```

```
print(category_count)
```

```
px.bar(category_count, x = 'count', y='Category', title='Category Count', orientation='h')
```

```
Category  count
0  Office Supplies  5360
1    Furniture     1880
2    Technology    1640
```

Category Count



Write your code here (7.3)

```
total = df_superstore_order['Category'].count()
```

```
Furniture = df_superstore_order[df_superstore_order['Category'] == 'Furniture']['Category'].count()
```

```
Technology = df_superstore_order[df_superstore_order['Category'] == 'Technology']['Category'].count()
```

```
Office_Supplies = df_superstore_order[df_superstore_order['Category'] == 'Office Supplies']['Category'].count()
```

```
print(f"Ratio of each Category: \nFurniture: {(Furniture/total * 100):.2f} % \nTechnology: {(Technology/total * 100):.2f} % \nOffice Supplies: {(Office_Supplies/total * 100):.2f} %")
```

```
Ratio of each Category:
Furniture: 21.17 %
Technology: 18.47 %
Office Supplies: 60.36 %
```

8. Update the type of all columns that contain dates to *datetime* and show information after an update.

write your code here

```
for col in df_superstore_order.columns:
```

```
    if 'Date' in col:
```

```
        df_superstore_order[col] = pd.to_datetime(df_superstore_order[col], format='%d/%m/%Y')
```

```
df_superstore_order.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 8880 entries, 0 to 8879
Data columns (total 21 columns):
#   Column      Non-Null Count  Dtype
---  ---
0   Row ID      8880 non-null  int64
1   Order ID    8880 non-null  object
2   Order Date  8880 non-null  datetime64[ns]
3   Ship Date   8880 non-null  datetime64[ns]
```

```

4 Ship Mode      8880 non-null object
5 Customer ID    8880 non-null object
6 Customer Name  8880 non-null object
7 Segment        8880 non-null object
8 Country        8880 non-null object
9 City           8880 non-null object
10 State          8880 non-null object
11 Postal Code    8880 non-null int64
12 Region         8880 non-null object
13 Product ID     8880 non-null object
14 Category       8880 non-null object
15 Sub-Category   8880 non-null object
16 Product Name   8880 non-null object
17 Sales          8880 non-null float64
18 Quantity       8880 non-null int64
19 Discount       8880 non-null float64
20 Profit         8880 non-null float64
dtypes: datetime64[ns](2), float64(3), int64(3), object(13)
memory usage: 1.4+ MB

```

```
df_superstore_order[[x for x in df_superstore_order.columns if 'Date' in x]].head()
```

	Order Date	Ship Date
0	2016-11-08	2016-11-11
1	2016-11-08	2016-11-11
2	2016-06-12	2016-06-16
3	2015-10-11	2015-10-18
4	2015-10-11	2015-10-18

9. Create a new column "Processing time day" to show number of days taken to ship an order and show your result in a dataframe format.

Hint: The duration starts as soon as the item has been ordered and ends once the order has successfully shipped.

```

# write your code here
df_superstore_order['Processing time day'] = (df_superstore_order['Ship Date'] - df_superstore_order['Order Date']).dt.days
# why .dt.days? because we want to get the difference in days
df_superstore_order['Processing time day'].head()

```

	Processing time day
0	3
1	3
2	4
3	7
4	7

10. Based on the result in 9.

- 10.1 How many orders are there that take more than 5 days to process?
- 10.2 Show the top 5 rows (expected output should contain these columns: Order ID, Order Date, Ship Date, Processing time day, Quantity)
- 10.3 Plot the histogram based on the column Quantity

Note: please create additional cells to answer 10.2 - 10.3

```

# Write your code here (10.1)
df_superstore_order[df_superstore_order['Processing time day'] > 5].shape
# so the answer is 1656

```

```
(1656, 22)
```

```

# write your code here (10.2)
df_superstore_order[['Order ID', 'Order Date', 'Ship Date', 'Processing time day', 'Quantity']].head(5)

```




	Order ID	Order Date	Ship Date	Processing time day	Quantity	
0	CA-2016-152156	2016-11-08	2016-11-11	3	2	
1	CA-2016-152156	2016-11-08	2016-11-11	3	3	
2	CA-2016-138688	2016-06-12	2016-06-16	4	2	
3	US-2015-108966	2015-10-11	2015-10-18	7	5	
4	US-2015-108966	2015-10-11	2015-10-18	7	2	

write your code here (10.3)

```
px.histogram(df_superstore_order['Quantity'], x='Quantity', title='Quantity Distribution')
```



Quantity Distribution



11. Total sales compare across different regions

- 11.1 Create a bar chart to visualize.

Write your code here (11.1)

```
sales_region = df_superstore_order.groupby(by = 'Region').agg(Total_Sales = ('Sales', 'sum')).reset_index()
sales_region
```



	Region	Total_Sales	
0	Central	434723.2222	
1	East	603648.0590	
2	South	361376.7910	
3	West	645943.7905	

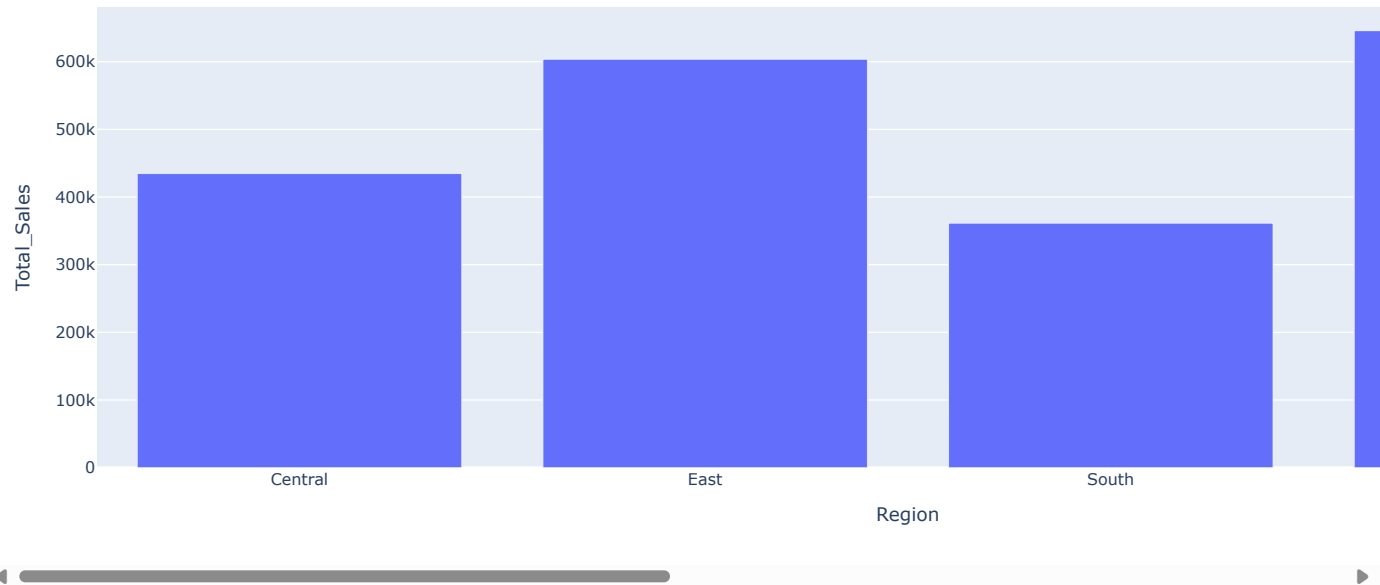
Next steps:

[Generate code with sales_region](#)
[View recommended plots](#)
[New interactive sheet](#)

```
px.bar(sales_region, x='Region', y='Total_Sales', title='Total Sales by Region')
```



Total Sales by Region



- 11.2 How do total sales compare across different regions? Explain in as much detail as possible.

Ans: The total sales in the West region is the highest with a total of 645943.79, followed by the East region with a total of 603648.06, then the Central region with a total of 434723.2222, and the lowest total sales in the South region with a total of 361376.79.

12.Which states have the highest number of returns? Use a horizontal bar chart.

Ans:

```
# Write your code here (12)
order_return = pd.merge(df_superstore_order, df_superstore_reutrnrn, on='Order ID', how='left')
order_return.fillna('No', inplace=True)
order_return.head()
```



	Row ID	Order ID	Order Date	Ship Date	Ship Mode	Customer ID	Customer Name	Segment	Country	City	...	Product ID	Category	Sub-Category	Product Name
0	1	CA-2016-152156	2016-11-08	2016-11-11	Second Class	CG-12520	Claire Gute	Consumer	United States	Henderson	...	FUR-BO-10001798	Furniture	Bookcases	Bush Somerset Collection Bookcase
1	2	CA-2016-152156	2016-11-08	2016-11-11	Second Class	CG-12520	Claire Gute	Consumer	United States	Henderson	...	FUR-CH-10000454	Furniture	Chairs	Hon Deluxe Fabric Upholstered Stacking Chairs ...
2	3	CA-2016-138688	2016-06-12	2016-06-16	Second Class	DV-13045	Darrin Van Huff	Corporate	United States	Los Angeles	...	OFF-LA-10000240	Office Supplies	Labels	Self-Adhesive Address Labels for Typewriters b...
3	4	US-2015-108966	2015-10-11	2015-10-18	Standard Class	SO-20335	Sean ODonnell	Consumer	United States	Fort Lauderdale	...	FUR-TA-10000577	Furniture	Tables	Bretford CR4500 Series Slim Rectangular Table
4	5	US-2015-108966	2015-10-11	2015-10-18	Standard Class	SO-20335	Sean ODonnell	Consumer	United States	Fort Lauderdale	...	OFF-ST-10000760	Office Supplies	Storage	Eldon Fold N Roll Cart System

5 rows x 23 columns


```
order_return.columns
```

```
Index(['Row ID', 'Order ID', 'Order Date', 'Ship Date', 'Ship Mode',
      'Customer ID', 'Customer Name', 'Segment', 'Country', 'City', 'State',
      'Postal Code', 'Region', 'Product ID', 'Category', 'Sub-Category',
      'Product Name', 'Sales', 'Quantity', 'Discount', 'Profit',
      'Processing time day', 'Returned'],
      dtype='object')
```


```
order_return_state = order_return.groupby(by = 'State').agg(Total_Returned = ('Returned', lambda x: (x == 'Yes').sum())).sort_values(by = 'Total_Returned')
```

```
#horizontal bar chart
```

```
px.bar(order_return_state, x='Total_Returned', y='State', title='Total Returned by State', orientation='h')
```





Total Returned by State



```
import matplotlib.pyplot as plt
```

```
plt.figure(figsize=(5, 20))
plt.barh(order_return_state['State'], order_return_state['Total_Returned'])
plt.xlabel('Total Returned')
plt.ylabel('State')
plt.title('Total Returned by State')
plt.show()
```





13. What is the correlation between numerical variables in the superstore_order dataset? Use a heatmap


Hint: Use seaborn to create a heatmap :)

```
# Write your code here (13)
import seaborn as sns

numeric = df_superstore_order.select_dtypes(include=['number']).columns
corr = df_superstore_order[numeric].corr()

sns.heatmap(corr, annot=True, cmap='coolwarm')
plt.title('Correlation Matrix')
```

↗ Text(0.5, 1.0, 'Correlation Matrix')



14. Create a USA State-Level Choropleth Map to visualize total sales per state.

- The darkest color represents the highest total sales.
- The lightest color represents the lowest total sales.
- Use a continuous gradient scale (e.g., dark blue to light blue, dark red to light red, or any custom gradient of your choice).

Hint: Use plotly.express

```
import plotly.express as px
state_sales = df_superstore_order.groupby(by = 'State').agg(Total_Sales = ('Sales', 'sum')).reset_index()
state_sales.head(10)
```

	State	Total_Sales
0	Alabama	19272.4900
1	Arizona	32990.1060
2	Arkansas	9427.4400
3	California	399195.4555
4	Colorado	31080.8060
5	Connecticut	12829.0570
6	Delaware	26379.4290
7	District of Columbia	2865.0200
8	Florida	84083.0880
9	Georgia	47940.1900

Next steps: [Generate code with state_sales](#) [View recommended plots](#) [New interactive sheet](#)


```
us_state_abbrev = {
    'Alabama': 'AL', 'Alaska': 'AK', 'Arizona': 'AZ', 'Arkansas': 'AR', 'California': 'CA',
    'Colorado': 'CO', 'Connecticut': 'CT', 'Delaware': 'DE', 'Florida': 'FL', 'Georgia': 'GA',
    'Hawaii': 'HI', 'Idaho': 'ID', 'Illinois': 'IL', 'Indiana': 'IN', 'Iowa': 'IA', 'Kansas': 'KS',
    'Kentucky': 'KY', 'Louisiana': 'LA', 'Maine': 'ME', 'Maryland': 'MD', 'Massachusetts': 'MA',
    'Michigan': 'MI', 'Minnesota': 'MN', 'Mississippi': 'MS', 'Missouri': 'MO', 'Montana': 'MT',
    'Nebraska': 'NE', 'Nevada': 'NV', 'New Hampshire': 'NH', 'New Jersey': 'NJ', 'New Mexico': 'NM',
    'New York': 'NY', 'North Carolina': 'NC', 'North Dakota': 'ND', 'Ohio': 'OH', 'Oklahoma': 'OK',
    'Oregon': 'OR', 'Pennsylvania': 'PA', 'Rhode Island': 'RI', 'South Carolina': 'SC',
    'South Dakota': 'SD', 'Tennessee': 'TN', 'Texas': 'TX', 'Utah': 'UT', 'Vermont': 'VT',
    'Virginia': 'VA', 'Washington': 'WA', 'West Virginia': 'WV', 'Wisconsin': 'WI', 'Wyoming': 'WY'
}

# Convert Full State Names to Abbreviations
state_sales['State'] = state_sales['State'].map(us_state_abbrev)

fig = px.choropleth(locations=state_sales['State'], locationmode='USA-states', color=state_sales['Total_Sales'], scope='usa', title='Total Sa
fig.show()
```



Total Sales per State



14.2 Answer the following questions:

1. Which state has the highest total sales?
2. How do sales anomalies affect the gradient color shading on the map?
3. If you change the color scale, does it impact readability? Why or why not?

Ans:


1. CA: California
2. Sales anomalies affect the gradient color shading on the map because it will make the map harder to read and interpret.
3. Yes, it will impact readability because the color scale will make the map more readable and easier to interpret. If the color scale is not chosen correctly, the map will be harder to interpret and understand.

15. Create a box plot to compare the different shipping modes based on total profit.

#Write your code here (15)

```
profit_shipMode = df_superstore_order[['Ship Mode', 'Profit']]

profit_shipMode.boxplot(by='Ship Mode', column='Profit', grid=False)
plt.show()
```



15.2 Which shipping mode has the highest median profit?

Ans: First Class

[BONUS 20 pts] Determine the percentage of customers who:

- B1) returned the product once
- B2) returned the product at least once
- B3) never returned the product
- Finally, Plot a comparison of B2 and B3

Note: please create additional cells to answer the above points

```
# Write your code here
Returned = order_return[order_return['Returned'] == 'Yes'].groupby('Customer ID').agg(return_count = ('Returned', 'count')).reset_index()
Returned_once = Returned[Returned['return_count'] == 1].shape[0]
Returned_atleast_once = Returned[Returned['return_count'] >= 1].shape[0]
Returned_never = order_return[order_return['Returned'] == 'No']['Customer ID'].unique().shape[0]
```

```
Returned_once_ratio = Returned_once / (Returned_once + Returned_atleast_once + Returned_never) * 100
Returned_atleast_once_ratio = Returned_atleast_once / (Returned_once + Returned_atleast_once + Returned_never) * 100
Returned_never_ratio = Returned_never / (Returned_once + Returned_atleast_once + Returned_never) * 100
```

```
Returned_once_ratio, Returned_atleast_once_ratio, Returned_never_ratio
```

```
(5.794392523364486, 20.747663551401867, 73.45794392523365)
```

```
# plot a comparison returned at least once, and never returned customers
px.pie(values=[Returned_atleast_once_ratio, Returned_never_ratio], names=['At least Once', 'Never'], title='Comparison of Returned at least
```



Comparison of Returned at least Once and Never Returned Customers

