

Lab 3: Data Preparation

CPE232 Data Models

✓ [1] Reviews on Pandas

1.1) Discover

- methods to explore and understand your DataFrame

```
import pandas as pd
```

```
df = pd.read_csv('nss15.csv')
```

```
# see the shape of the dataframe
print(df.shape)
```

```
↗ (334839, 12)
```

```
# seeing the summary of the dataframe
print(df.info())
```

```
↗ <class 'pandas.core.frame.DataFrame'>
RangeIndex: 334839 entries, 0 to 334838
Data columns (total 12 columns):
#   Column              Non-Null Count  Dtype
---  -
0   caseNumber          334839 non-null  int64
1   treatmentDate       334839 non-null  object
2   statWeight          334839 non-null  float64
3   stratum             334839 non-null  object
4   age                 334839 non-null  int64
5   sex                 334837 non-null  object
6   race                205014 non-null  object
7   diagnosis           334839 non-null  int64
8   bodyPart            334839 non-null  int64
9   disposition         334839 non-null  int64
10  location            334839 non-null  int64
11  product             334839 non-null  int64
dtypes: float64(1), int64(7), object(4)
memory usage: 30.7+ MB
None
```

```
# seeing the stats of the column in dataframe
print(df.describe())
```

```
↗
```

	caseNumber	statWeight	age	diagnosis \
count	3.348390e+05	334839.000000	334839.000000	334839.000000
mean	1.510271e+08	39.343028	31.385451	60.154591
std	1.720330e+06	34.142933	26.105098	6.170699
min	1.501032e+08	4.965500	0.000000	41.000000
25%	1.504405e+08	15.059100	10.000000	57.000000
50%	1.507358e+08	15.776200	23.000000	59.000000
75%	1.510231e+08	74.881300	51.000000	64.000000
max	1.603418e+08	97.923900	107.000000	74.000000

	bodyPart	disposition	location	product
count	334839.000000	334839.000000	334839.000000	334839.000000
mean	64.374192	1.307930	2.485451	2098.900854
std	24.002331	0.977627	3.217617	1332.222670
min	0.000000	1.000000	0.000000	106.000000
25%	35.000000	1.000000	0.000000	1211.000000
50%	75.000000	1.000000	1.000000	1807.000000
75%	82.000000	1.000000	5.000000	3265.000000
max	94.000000	9.000000	9.000000	5555.000000

```
# seeing the first 5 rows of the dataframe
print(df.head())
```

```
↗
```

	caseNumber	treatmentDate	statWeight	stratum	age	sex	race \
0	150733174	7/11/2015	15.7762	V	5	Male	NaN
1	150734723	7/6/2015	83.2157	S	36	Male	White
2	150817487	8/2/2015	74.8813	L	20	Female	NaN
3	150717776	6/26/2015	15.7762	V	61	Male	NaN
4	150721694	7/4/2015	74.8813	L	88	Female	Other

	diagnosis	bodyPart	disposition	location	product
0	57	33	1	9	1267
1	57	34	1	1	1439
2	71	94	1	0	3274
3	71	35	1	0	611
4	62	75	1	0	1893

```
# seeing the last 5 rows of the dataframe
print(df.tail())
```

```
↵
```

	caseNumber	treatmentDate	statWeight	stratum	age	sex	race	\
334834	150739278	5/31/2015	15.0591	V	7	Male	NaN	
334835	150733393	7/11/2015	5.6748	C	3	Female	Black	
334836	150819286	7/24/2015	15.7762	V	38	Male	NaN	
334837	150823002	8/8/2015	97.9239	M	38	Female	White	
334838	150723074	6/20/2015	49.2646	M	5	Female	White	

	diagnosis	bodyPart	disposition	location	product
334834	59	76	1	1	1864
334835	68	85	1	0	1931
334836	71	79	1	0	3250
334837	59	82	1	1	464
334838	57	34	1	9	3273

```
# seeing the list of columns in the dataframe
print(df.columns)
```

```
↵ Index(['caseNumber', 'treatmentDate', 'statWeight', 'stratum', 'age', 'sex',
        'race', 'diagnosis', 'bodyPart', 'disposition', 'location', 'product'],
        dtype='object')
```

1.2) Selecting variables

- select specific columns from the DataFrame to create a new DataFrame with only those columns

```
df['age']
```

```
↵
```

0	5
1	36
2	20
3	61
4	88
	..
334834	7
334835	3
334836	38
334837	38
334838	5

Name: age, Length: 334839, dtype: int64

```
df['age'].head()
```

```
↵
```

0	5
1	36
2	20
3	61
4	88

Name: age, dtype: int64

```
df[['caseNumber', 'age']]
```



	caseNumber	age
0	150733174	5
1	150734723	36
2	150817487	20
3	150717776	61
4	150721694	88
...
334834	150739278	7
334835	150733393	3
334836	150819286	38
334837	150823002	38
334838	150723074	5

334839 rows × 2 columns

```
# select columns based on the data type
df.select_dtypes(include=['number'])
```



	caseNumber	statWeight	age	diagnosis	bodyPart	disposition	location	product
0	150733174	15.7762	5	57	33	1	9	1267
1	150734723	83.2157	36	57	34	1	1	1439
2	150817487	74.8813	20	71	94	1	0	3274
3	150717776	15.7762	61	71	35	1	0	611
4	150721694	74.8813	88	62	75	1	0	1893
...
334834	150739278	15.0591	7	59	76	1	1	1864
334835	150733393	5.6748	3	68	85	1	0	1931
334836	150819286	15.7762	38	71	79	1	0	3250
334837	150823002	97.9239	38	59	82	1	1	464
334838	150723074	49.2646	5	57	34	1	9	3273

334839 rows × 8 columns

```
# select row by .loc
df.loc[0]
```



```
caseNumber      150733174
treatmentDate    7/11/2015
statWeight       15.7762
stratum          V
age              5
sex              Male
race             NaN
diagnosis        57
bodyPart         33
disposition      1
location         9
product         1267
Name: 0, dtype: object
```

```
# select column by .loc
df.loc[:, 'treatmentDate': 'diagnosis']
```



	treatmentDate	statWeight	stratum	age	sex	race	diagnosis
0	7/11/2015	15.7762	V	5	Male	NaN	57
1	7/6/2015	83.2157	S	36	Male	White	57
2	8/2/2015	74.8813	L	20	Female	NaN	71
3	6/26/2015	15.7762	V	61	Male	NaN	71
4	7/4/2015	74.8813	L	88	Female	Other	62
5	7/2/2015	5.6748	C	1	Female	White	71
6	6/8/2015	15.7762	V	25	Male	Black	51

```
df.loc[df['age']>80, ['treatmentDate', 'age']]
```



	treatmentDate	age
4	7/4/2015	88
8	7/16/2015	98
39	5/3/2015	88
46	4/15/2015	91
63	1/12/2015	97
...
334701	4/27/2015	86
334784	7/7/2015	82
334785	7/11/2015	86
334815	10/28/2015	85
334819	1/13/2015	85

20422 rows × 2 columns

```
# select row by .iloc
df.iloc[0:5]
```



	caseNumber	treatmentDate	statWeight	stratum	age	sex	race	diagnosis	bodyPart	disposition	location	product
0	150733174	7/11/2015	15.7762	V	5	Male	NaN	57	33	1	9	1267
1	150734723	7/6/2015	83.2157	S	36	Male	White	57	34	1	1	1439
2	150817487	8/2/2015	74.8813	L	20	Female	NaN	71	94	1	0	3274
3	150717776	6/26/2015	15.7762	V	61	Male	NaN	71	35	1	0	611
4	150721694	7/4/2015	74.8813	L	88	Female	Other	62	75	1	0	1893

```
# select column by .iloc
df.iloc[:,[0,1,2,3,4]]
```



	caseNumber	treatmentDate	statWeight	stratum	age
0	150733174	7/11/2015	15.7762	V	5
1	150734723	7/6/2015	83.2157	S	36
2	150817487	8/2/2015	74.8813	L	20
3	150717776	6/26/2015	15.7762	V	61
4	150721694	7/4/2015	74.8813	L	88
...
334834	150739278	5/31/2015	15.0591	V	7
334835	150733393	7/11/2015	5.6748	C	3
334836	150819286	7/24/2015	15.7762	V	38
334837	150823002	8/8/2015	97.9239	M	38
334838	150723074	6/20/2015	49.2646	M	5

334839 rows × 5 columns

1.3) Filtering the data

```
# filter rows based on the condition
df[df['age'] > 50]
```

	caseNumber	treatmentDate	statWeight	stratum	age	sex	race	diagnosis	bodyPart	disposition	location	product	
	3	150717776	6/26/2015	15.7762	V	61	Male	NaN	71	35	1	0	611
	4	150721694	7/4/2015	74.8813	L	88	Female	Other	62	75	1	0	1893
	7	150704114	6/14/2015	83.2157	S	53	Male	White	57	30	1	0	5040
	8	150736558	7/16/2015	83.2157	S	98	Male	Black	59	76	1	1	1807
	16	150901411	8/27/2015	83.2157	S	65	Female	White	59	83	1	1	1817
...
334811	150702215	6/27/2015	15.7762	V	51	Female	NaN	53	83	1	1	1426	
334815	151100368	10/28/2015	83.2157	S	85	Female	NaN	57	80	4	1	1807	
334819	150528367	1/13/2015	49.2646	M	85	Female	NaN	57	79	5	1	676	
334826	150648619	6/17/2015	15.7762	V	52	Female	White	64	30	1	1	1842	
334829	150633526	4/4/2015	49.2646	M	51	Female	NaN	56	92	1	1	1616	

85235 rows × 12 columns

```
# filter coloum based on column name
df.filter(like='age')
```




	age
0	5
1	36
2	20
3	61
4	88
...	...
334834	7
334835	3
334836	38
334837	38
334838	5

334839 rows × 1 columns

1.4) Sorting

- Sort the DataFrame by its index based on column


```
# sort the dataframe based on column name and ascending order
df.sort_values(by='statWeight', ascending=False)
```



	caseNumber	treatmentDate	statWeight	stratum	age	sex	race	diagnosis	bodyPart	disposition	location	product
275174	150343700	3/9/2015	97.9239	M	48	Male	NaN	57	93	1	1	281
36	151029422	10/6/2015	97.9239	M	37	Male	White	64	35	1	0	1267
334806	150612491	5/29/2015	97.9239	M	18	Female	White	59	92	1	1	845
334810	150725804	7/8/2015	97.9239	M	33	Female	Black	71	94	1	0	1616
275161	150450816	4/13/2015	97.9239	M	24	Male	White	71	37	1	1	3286
...
44011	160222258	12/29/2015	4.9655	C	2	Female	Other	71	92	1	1	1893
325320	151213065	11/29/2015	4.9655	C	16	Female	White	62	75	1	8	3254
43891	160113865	12/28/2015	4.9655	C	4	Male	White	59	76	1	1	1842
43628	151130111	11/9/2015	4.9655	C	13	Male	Black	53	33	1	0	5011
43523	151139237	11/16/2015	4.9655	C	2	Female	Black	57	80	1	0	679

334839 rows × 12 columns

```
# sort the index of the dataframe
df.sort_index()
```




	caseNumber	treatmentDate	statWeight	stratum	age	sex	race	diagnosis	bodyPart	disposition	location	product
0	150733174	7/11/2015	15.7762	V	5	Male	NaN	57	33	1	9	1267
1	150734723	7/6/2015	83.2157	S	36	Male	White	57	34	1	1	1439
2	150817487	8/2/2015	74.8813	L	20	Female	NaN	71	94	1	0	3274
3	150717776	6/26/2015	15.7762	V	61	Male	NaN	71	35	1	0	611
4	150721694	7/4/2015	74.8813	L	88	Female	Other	62	75	1	0	1893
...
334834	150739278	5/31/2015	15.0591	V	7	Male	NaN	59	76	1	1	1864
334835	150733393	7/11/2015	5.6748	C	3	Female	Black	68	85	1	0	1931
334836	150819286	7/24/2015	15.7762	V	38	Male	NaN	71	79	1	0	3250
334837	150823002	8/8/2015	97.9239	M	38	Female	White	59	82	1	1	464
334838	150723074	6/20/2015	49.2646	M	5	Female	White	57	34	1	9	3273

334839 rows × 12 columns

1.5) Add/Remove

- This section shows how to manipulate the DataFrame's structure


```
# Dropping the column
df.drop(columns=['disposition'])
```



	caseNumber	treatmentDate	statWeight	stratum	age	sex	race	diagnosis	bodyPart	location	product
0	150733174	7/11/2015	15.7762	V	5	Male	NaN	57	33	9	1267
1	150734723	7/6/2015	83.2157	S	36	Male	White	57	34	1	1439
2	150817487	8/2/2015	74.8813	L	20	Female	NaN	71	94	0	3274
3	150717776	6/26/2015	15.7762	V	61	Male	NaN	71	35	0	611
4	150721694	7/4/2015	74.8813	L	88	Female	Other	62	75	0	1893
...
334834	150739278	5/31/2015	15.0591	V	7	Male	NaN	59	76	1	1864
334835	150733393	7/11/2015	5.6748	C	3	Female	Black	68	85	0	1931
334836	150819286	7/24/2015	15.7762	V	38	Male	NaN	71	79	0	3250
334837	150823002	8/8/2015	97.9239	M	38	Female	White	59	82	1	464
334838	150723074	6/20/2015	49.2646	M	5	Female	White	57	34	9	3273

334839 rows × 11 columns

```
# Adding column and create into a new column
df.assign(new_column=df['diagnosis'] + df['bodyPart'])
```



	caseNumber	treatmentDate	statWeight	stratum	age	sex	race	diagnosis	bodyPart	disposition	location	product	new_
0	150733174	7/11/2015	15.7762	V	5	Male	NaN	57	33	1	9	1267	
1	150734723	7/6/2015	83.2157	S	36	Male	White	57	34	1	1	1439	
2	150817487	8/2/2015	74.8813	L	20	Female	NaN	71	94	1	0	3274	
3	150717776	6/26/2015	15.7762	V	61	Male	NaN	71	35	1	0	611	
4	150721694	7/4/2015	74.8813	L	88	Female	Other	62	75	1	0	1893	
...
334834	150739278	5/31/2015	15.0591	V	7	Male	NaN	59	76	1	1	1864	
334835	150733393	7/11/2015	5.6748	C	3	Female	Black	68	85	1	0	1931	
334836	150819286	7/24/2015	15.7762	V	38	Male	NaN	71	79	1	0	3250	
334837	150823002	8/8/2015	97.9239	M	38	Female	White	59	82	1	1	464	
334838	150723074	6/20/2015	49.2646	M	5	Female	White	57	34	1	9	3273	


334839 rows × 13 columns

```
# Removing the column and assigning it to a new variable
ages = df.pop('age')
```

1.6) Clean missing

- to remove rows with missing values or replace missing values with a specified value


```
# replacing the missing values with a specified value
df.fillna(value=0)
```



	caseNumber	treatmentDate	statWeight	stratum	sex	race	diagnosis	bodyPart	disposition	location	product
0	150733174	7/11/2015	15.7762	V	Male	0	57	33	1	9	1267
1	150734723	7/6/2015	83.2157	S	Male	White	57	34	1	1	1439
2	150817487	8/2/2015	74.8813	L	Female	0	71	94	1	0	3274
3	150717776	6/26/2015	15.7762	V	Male	0	71	35	1	0	611
4	150721694	7/4/2015	74.8813	L	Female	Other	62	75	1	0	1893
...
334834	150739278	5/31/2015	15.0591	V	Male	0	59	76	1	1	1864
334835	150733393	7/11/2015	5.6748	C	Female	Black	68	85	1	0	1931
334836	150819286	7/24/2015	15.7762	V	Male	0	71	79	1	0	3250
334837	150823002	8/8/2015	97.9239	M	Female	White	59	82	1	1	464
334838	150723074	6/20/2015	49.2646	M	Female	White	57	34	1	9	3273

334839 rows × 11 columns

```
# Remove the rows with missing values
df.dropna()
```



	caseNumber	treatmentDate	statWeight	stratum	sex	race	diagnosis	bodyPart	disposition	location	product
1	150734723	7/6/2015	83.2157	S	Male	White	57	34	1	1	1439
4	150721694	7/4/2015	74.8813	L	Female	Other	62	75	1	0	1893
5	150721815	7/2/2015	5.6748	C	Female	White	71	76	1	1	1715
6	150713483	6/8/2015	15.7762	V	Male	Black	51	33	4	9	1138
7	150704114	6/14/2015	83.2157	S	Male	White	57	30	1	0	5040
...
334830	150628863	6/8/2015	15.7762	V	Female	White	64	79	1	1	1522
334831	150607637	5/22/2015	5.6748	C	Female	Black	59	94	1	0	1616
334835	150733393	7/11/2015	5.6748	C	Female	Black	68	85	1	0	1931
334837	150823002	8/8/2015	97.9239	M	Female	White	59	82	1	1	464
334838	150723074	6/20/2015	49.2646	M	Female	White	57	34	1	9	3273


205014 rows × 11 columns

✓ [2] Data Cleaning and Preparation

✓ .isnull, .dropna, .fillna


2.1) checking

```
df.columns
```



```
Index(['caseNumber', 'treatmentDate', 'statWeight', 'stratum', 'sex', 'race',
      'diagnosis', 'bodyPart', 'disposition', 'location', 'product'],
      dtype='object')
```

```
# isnull checking
df.isnull().sum()
```



```
caseNumber      0
treatmentDate   0
statWeight      0
```

```

stratum      0
sex          2
race        129825
diagnosis    0
bodyPart     0
disposition  0
location     0
product      0
dtype: int64

```

```

# percentage of missing values for the race
df.race.isnull().sum()/df.shape[0]*100

```

```
np.float64(38.772365226272925)
```

```
df.shape[0]
```

```
334839
```

2.2) Drop column

```

# remove column by using
df = df.drop(columns=['race'])

```

```
df.head()
```

```

caseNumber  treatmentDate  statWeight  stratum  sex  diagnosis  bodyPart  disposition  location  product
0  150733174      7/11/2015    15.7762      V   Male      57      33           1           9      1267
1  150734723      7/6/2015    83.2157      S   Male      57      34           1           1      1439
2  150817487      8/2/2015    74.8813      L  Female      71      94           1           0      3274
3  150717776      6/26/2015    15.7762      V   Male      71      35           1           0       611
4  150721694      7/4/2015    74.8813      L  Female      62      75           1           0      1893

```

2.3) Data imputation

```

# fillna
df['age'] = df['age'].fillna(df['age'].median())

```




```

-----
KeyError                                Traceback (most recent call last)
File c:\Users\adiso\AppData\Local\Programs\Python\Python313\Lib\site-packages\pandas\core\indexes\base.py:3805, in
Index.get_loc(self, key)
    3804 try:
-> 3805     return self._engine.get_loc(casted_key)
    3806 except KeyError as err:

File index.pyx:167, in pandas._libs.index.IndexEngine.get_loc()

File index.pyx:196, in pandas._libs.index.IndexEngine.get_loc()

File pandas\_libs\hashtable_class_helper.pxi:7081, in pandas._libs.hashtable.PyObjectHashTable.get_item()

File pandas\_libs\hashtable_class_helper.pxi:7089, in pandas._libs.hashtable.PyObjectHashTable.get_item()

KeyError: 'age'

The above exception was the direct cause of the following exception:

KeyError                                Traceback (most recent call last)
Cell In[32], line 2
      1 # fillna
----> 2 df['age'] = df['age'].fillna(df['age'].median())

File c:\Users\adiso\AppData\Local\Programs\Python\Python313\Lib\site-packages\pandas\core\frame.py:4102, in
DataFrame._getitem_(self, key)
    4100 if self.columns.nlevels > 1:
    4101     return self._getitem_multilevel(key)
-> 4102 indexer = self.columns.get_loc(key)
    4103 if is_integer(indexer):
    4104     indexer = [indexer]

File c:\Users\adiso\AppData\Local\Programs\Python\Python313\Lib\site-packages\pandas\core\indexes\base.py:3812, in
Index.get_loc(self, key)
    3807 if isinstance(casted_key, slice) or (
    3808     isinstance(casted_key, abc.Iterable)
    3809     and any(isinstance(x, slice) for x in casted_key)
    3810 ):
    3811     raise InvalidIndexError(key)
-> 3812 raise KeyError(key) from err
    3813 except TypeError:
    3814     # If we have a Listlike key, _check_indexing_error will raise
    3815     # InvalidIndexError. Otherwise we fall through and re-raise
    3816     # the TypeError.
    3817     self._check_indexing_error(key)

```

[Q1] From the above cell, Why it showing an error?

Ans: The error message "KeyError: 'age'" indicates that the DataFrame does not contain a column named 'age'. Upon checking with `df.columns`, the available columns are:

`Index(['caseNumber', 'treatmentDate', 'statWeight', 'stratum', 'sex', 'race', 'diagnosis', 'bodyPart', 'disposition', 'location', 'product'], dtype='object')`

This confirms that the 'age' column is missing from the DataFrame.

[Q2] Fix the error from Q1 problem.

[Q2]

hint: see the cell that run `df.pop()`

`df["age"] = ages`

fillna again

`df['age'] = df['age'].fillna(df['age'].median())`

`df.head()`



	caseNumber	treatmentDate	statWeight	stratum	sex	diagnosis	bodyPart	disposition	location	product	age	Year	Month
0	150733174	11/07/2015	15.7762	V	Male	57	33	1	9	1267	5	2015	7
1	150734723	06/07/2015	83.2157	S	Male	57	34	1	1	1439	36	2015	7
2	150817487	02/08/2015	74.8813	L	Female	71	94	1	0	3274	20	2015	8
3	150717776	26/06/2015	15.7762	V	Male	71	35	1	0	611	61	2015	6
4	150721694	04/07/2015	74.8813	L	Female	62	75	1	0	1893	88	2015	7

2.4) Drop row that have missing value

```
# remove column by using .dropna()
df = df.dropna()
```

```
df.isnull().sum()
```

```
caseNumber      0
treatmentDate   0
statWeight      0
stratum         0
sex             0
diagnosis       0
bodyPart        0
disposition     0
location        0
product         0
age             0
dtype: int64
```

▼ Datetime

2.5) Working with the datetime format

```
df["treatmentDate"] = pd.to_datetime(df["treatmentDate"], format="%m/%d/%Y")
```

```
C:\Users\adiso\AppData\Local\Temp\ipykernel_17212\3208943844.py:1: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus

```
df["treatmentDate"] = pd.to_datetime(df["treatmentDate"], format="%m/%d/%Y")
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 334837 entries, 0 to 334838
Data columns (total 11 columns):
#   Column          Non-Null Count  Dtype
---  -
0   caseNumber      334837 non-null  int64
1   treatmentDate   334837 non-null  datetime64[ns]
2   statWeight      334837 non-null  float64
3   stratum         334837 non-null  object
4   sex             334837 non-null  object
5   diagnosis       334837 non-null  int64
6   bodyPart        334837 non-null  int64
7   disposition     334837 non-null  int64
8   location        334837 non-null  int64
9   product         334837 non-null  int64
10  age             334837 non-null  int64
dtypes: datetime64[ns](1), float64(1), int64(7), object(2)
memory usage: 30.7+ MB
```

```
df['Year'] = df['treatmentDate'].dt.year
```

```
C:\Users\adiso\AppData\Local\Temp\ipykernel_17212\1686165144.py:1: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus

```
df['Year'] = df['treatmentDate'].dt.year
```

```
df['Month'] = df['treatmentDate'].dt.month
```

```
C:\Users\adiso\AppData\Local\Temp\ipykernel_17212\404848564.py:1: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus

```
df['Month'] = df['treatmentDate'].dt.month
```

```
df.head()
```

	caseNumber	treatmentDate	statWeight	stratum	sex	diagnosis	bodyPart	disposition	location	product	age	Year	Month
0	150733174	2015-07-11	15.7762	V	Male	57	33	1	9	1267	0	2015	7
1	150734723	2015-07-06	83.2157	S	Male	57	34	1	1	1439	0	2015	7
2	150817487	2015-08-02	74.8813	L	Female	71	94	1	0	3274	0	2015	8
3	150717776	2015-06-26	15.7762	V	Male	71	35	1	0	611	0	2015	6
4	150721694	2015-07-04	74.8813	L	Female	62	75	1	0	1893	0	2015	7

[Q3] Can you change the format to DD/MM/YYYY? Show your work.

write your code here

```
df['treatmentDate'] = pd.to_datetime(df['treatmentDate'])
```

```
df['treatmentDate'] = df['treatmentDate'].dt.strftime('%d/%m/%Y')
```

```
print(df.head())
```

C:\Users\adiso\AppData\Local\Temp\ipykernel_17212\563021845.py:2: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus

```
df['treatmentDate'] = pd.to_datetime(df['treatmentDate'])
caseNumber treatmentDate statWeight stratum sex diagnosis bodyPart \
0 150733174 11/07/2015 15.7762 V Male 57 33
1 150734723 06/07/2015 83.2157 S Male 57 34
2 150817487 02/08/2015 74.8813 L Female 71 94
3 150717776 26/06/2015 15.7762 V Male 71 35
4 150721694 04/07/2015 74.8813 L Female 62 75
```

```
disposition location product age Year Month
0 1 9 1267 0 2015 7
1 1 1 1439 0 2015 7
2 1 0 3274 0 2015 8
3 1 0 611 0 2015 6
4 1 0 1893 0 2015 7
```

C:\Users\adiso\AppData\Local\Temp\ipykernel_17212\563021845.py:4: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus
df['treatmentDate'] = df['treatmentDate'].dt.strftime('%d/%m/%Y')

✓ Combine Dataframe by .merge and .concat

2.6 Merge

```
import pandas as pd
```

```
superstore_order = pd.read_csv('superstore_order.csv')
superstore_people = pd.read_csv('superstore_people.csv')
superstore_return = pd.read_csv('superstore_return.csv')
```

```
superstore_order.merge(superstore_return[superstore_return["Returned"]=="Yes"],
on="Order ID",
how="inner")\
[["Customer ID", "Returned"]]\
.drop_duplicates()
```



	Customer ID	Returned
0	ZD-21925	Yes
3	TB-21055	Yes
10	JS-15685	Yes
13	LC-16885	Yes
20	BS-11755	Yes
...
688	ED-13885	Yes
689	TS-21205	Yes
696	MF-17665	Yes
702	SH-19975	Yes
705	RB-19435	Yes

222 rows × 2 columns

[Q4] What does the argument `how="inner"` do?


Ans: argument `how="inner"` is to ensures that only rows with matching "Order ID" values in both DataFrames are retained and to filtering out unmatched records.

[Q5] In your opinion, what information that the result above conveys?

Ans: List of Customers Who Returned Orders and also provide Customer ID column shows which customers. the infomation are removing Duplicates it can use to analyze in to find patterns.

More merging...

```
superstore_order.merge(superstore_return,  
                        on="Order ID" ,  
                        how="inner")
```



	Row ID	Order ID	Order Date	Ship Date	Ship Mode	Customer ID	Customer Name	Segment	Country	City	...	Region	Product ID	Categor
0	19	CA-2014-143336	27/08/2014	01/09/2014	Second Class	ZD-21925	Zuschuss Donatelli	Consumer	United States	San Francisco	...	West	OFF-AR-10003056	Office Supplies
1	20	CA-2014-143336	27/08/2014	01/09/2014	Second Class	ZD-21925	Zuschuss Donatelli	Consumer	United States	San Francisco	...	West	TEC-PH-10001949	Technology
2	21	CA-2014-143336	27/08/2014	01/09/2014	Second Class	ZD-21925	Zuschuss Donatelli	Consumer	United States	San Francisco	...	West	OFF-BI-10002215	Office Supplies
3	56	CA-2016-111682	17/06/2016	18/06/2016	First Class	TB-21055	Ted Butterfield	Consumer	United States	Troy	...	East	OFF-ST-10000604	Office Supplies
4	57	CA-2016-111682	17/06/2016	18/06/2016	First Class	TB-21055	Ted Butterfield	Consumer	United States	Troy	...	East	OFF-PA-10001569	Office Supplies
...
702	8870	CA-2017-101805	01/12/2017	06/12/2017	Standard Class	SH-19975	Sally Hughsby	Corporate	United States	Seattle	...	West	OFF-BI-10002003	Office Supplies
703	8871	CA-2017-101805	01/12/2017	06/12/2017	Standard Class	SH-19975	Sally Hughsby	Corporate	United States	Seattle	...	West	FUR-FU-10000023	Furniture
704	8872	CA-2017-101805	01/12/2017	06/12/2017	Standard Class	SH-19975	Sally Hughsby	Corporate	United States	Seattle	...	West	OFF-ST-10002756	Office Supplies
705	8873	US-2014-105137	10/10/2014	10/10/2014	Same Day	RB-19435	Richard Bierner	Consumer	United States	Columbus	...	East	TEC-MA-10002694	Technology
706	8874	US-2014-105137	10/10/2014	10/10/2014	Same Day	RB-19435	Richard Bierner	Consumer	United States	Columbus	...	East	OFF-BI-10002429	Office Supplies

707 rows × 22 columns

2.7) Concatenate

```
pd.concat([superstore_order, superstore_people], axis=1, join='inner')
```



	Row ID	Order ID	Order Date	Ship Date	Ship Mode	Customer ID	Customer Name	Segment	Country	City	...	Product ID	Category	Sub-Category
0	1	CA-2016-152156	08/11/2016	11/11/2016	Second Class	CG-12520	Claire Gute	Consumer	United States	Henderson	...	FUR-BO-10001798	Furniture	Bookcases
1	2	CA-2016-152156	08/11/2016	11/11/2016	Second Class	CG-12520	Claire Gute	Consumer	United States	Henderson	...	FUR-CH-10000454	Furniture	Chairs
2	3	CA-2016-138688	12/06/2016	16/06/2016	Second Class	DV-13045	Darrin Van Huff	Corporate	United States	Los Angeles	...	OFF-LA-10000240	Office Supplies	Labels
3	4	US-2015-108966	11/10/2015	18/10/2015	Standard Class	SO-20335	Sean ODonnell	Consumer	United States	Fort Lauderdale	...	FUR-TA-10000577	Furniture	Tables


4 rows × 23 columns

[Q6] What is the difference between inner and outer on parameter join in pd.concat?

Ans: inner: Keeps only indices that appear in the input DataFrames. outer: Keeps all indices that appear in any of the input DataFrames then using NULL to fill missing values.

▼ Groupby

```
superstore_order.groupby(['Segment', 'Ship Mode'])[['Sales', 'Quantity', 'Discount', 'Profit']].sum()
```




			Sales	Quantity	Discount	Profit
Segment	Ship Mode					
Consumer	First Class		138594.9328	2455	110.29	18953.7264
	Same Day		53660.6340	1001	43.85	8555.7193
	Second Class		203605.6822	3489	127.29	24701.9148
	Standard Class		627061.3262	10430	443.05	68864.9892
Corporate	First Class		97720.1209	1670	73.07	12660.2526
	Same Day		41716.5550	366	14.50	1120.9222
	Second Class		130759.9288	2027	71.47	15582.1762
	Standard Class		359359.2109	6203	262.82	49832.6780
Home Office	First Class		76743.8674	924	39.82	11829.8821
	Same Day		20968.5170	343	12.50	3909.3442
	Second Class		77175.1080	1148	37.80	12785.8953
	Standard Class		218325.9795	3595	142.14	27298.5786

[Q7] Describe an information that the result above conveys?

Ans: Standard Class shipping is the highest volume shipping method. The profit columns show positive values on all segments and shipping modes

```
superstore_order["Profit Ratio"] = superstore_order["Profit"]/superstore_order["Sales"]
```

```
superstore_order.groupby(["Category", "Sub-Category"]).agg(mean_profit_ratio = ("Profit Ratio", "mean"))
```



		mean_profit_ratio
Category	Sub-Category	
Furniture	Bookcases	-0.127756
	Chairs	0.045028
	Furnishings	0.140782
	Tables	-0.147916
Office Supplies	Appliances	-0.145513
	Art	0.251678
	Binders	-0.191641
	Envelopes	0.421913
	Fasteners	0.301157
	Labels	0.429984
	Paper	0.425586
	Storage	0.092382
	Supplies	0.104970
Technology	Accessories	0.219012
	Copiers	0.317826
	Machines	-0.059535
	Phones	0.118926

[Q8] Describe an information that the result above conveys?

Ans: Some sub-categories have negative profit ratios, indicating losses: Bookcases (-0.12756) Tables (-0.14792) Appliances (-0.14551) Binders (-0.19164) Machines (-0.05954)

Technology category mostly shows positive profit ratios:

Copiers having the highest (0.31783)

Only Machines showing a slight loss (-0.05954)

Office Supplies has several highly profitable sub-categories:

Labels (0.42998)

Paper (0.42559)

Envelopes (0.42191)


Fasteners (0.30116)

Art (0.25168)

✓ Pivot and Melt

Pivot

```
superstore_order.pivot_table(index="State", columns="Ship Mode", values="Order ID", aggfunc="count").fillna(0).head(10)
```



	Ship Mode	First Class	Same Day	Second Class	Standard Class
State					
Alabama		9.0	1.0	18.0	30.0
Arizona		42.0	15.0	22.0	123.0
Arkansas		10.0	2.0	8.0	35.0
California		302.0	106.0	346.0	1000.0
Colorado		43.0	5.0	32.0	95.0
Connecticut		19.0	8.0	11.0	39.0
Delaware		16.0	2.0	13.0	55.0
District of Columbia		0.0	0.0	3.0	7.0
Florida		47.0	25.0	57.0	210.0
Georgia		19.0	15.0	31.0	108.0

```
pivot_table_result = superstore_order.pivot_table(index="State", columns="Ship Mode", values="Order ID", aggfunc="count").fillna(0)
print(pivot_table_result)
```

Ship Mode	First Class	Same Day	Second Class	Standard Class
State				
Alabama	9.0	1.0	18.0	30.0
Arizona	42.0	15.0	22.0	123.0
Arkansas	10.0	2.0	8.0	35.0
California	302.0	106.0	346.0	1000.0
Colorado	43.0	5.0	32.0	95.0
Connecticut	19.0	8.0	11.0	39.0
Delaware	16.0	2.0	13.0	55.0
District of Columbia	0.0	0.0	3.0	7.0
Florida	47.0	25.0	57.0	210.0
Georgia	19.0	15.0	31.0	108.0
Idaho	3.0	0.0	2.0	13.0
Illinois	58.0	24.0	96.0	249.0
Indiana	13.0	3.0	30.0	79.0
Iowa	1.0	1.0	4.0	17.0
Kansas	6.0	1.0	2.0	15.0
Kentucky	12.0	5.0	49.0	62.0
Louisiana	7.0	2.0	14.0	15.0
Maine	0.0	0.0	0.0	5.0
Maryland	18.0	7.0	12.0	63.0
Massachusetts	14.0	4.0	35.0	71.0
Michigan	20.0	16.0	43.0	151.0
Minnesota	9.0	4.0	13.0	59.0
Mississippi	3.0	4.0	7.0	36.0
Missouri	7.0	2.0	20.0	24.0
Montana	1.0	1.0	0.0	13.0
Nebraska	6.0	3.0	6.0	20.0
Nevada	4.0	1.0	12.0	17.0
New Hampshire	2.0	0.0	10.0	13.0
New Jersey	5.0	1.0	20.0	87.0
New Mexico	1.0	0.0	9.0	22.0
New York	155.0	57.0	183.0	606.0
North Carolina	36.0	14.0	40.0	139.0
North Dakota	0.0	0.0	5.0	2.0
Ohio	66.0	47.0	84.0	199.0
Oklahoma	5.0	6.0	7.0	44.0
Oregon	20.0	0.0	15.0	81.0
Pennsylvania	103.0	9.0	78.0	341.0
Rhode Island	16.0	0.0	21.0	16.0
South Carolina	3.0	5.0	18.0	16.0
South Dakota	2.0	0.0	0.0	9.0
Tennessee	21.0	2.0	24.0	118.0
Texas	125.0	37.0	161.0	537.0
Utah	4.0	2.0	19.0	28.0
Vermont	0.0	0.0	1.0	2.0
Virginia	39.0	4.0	33.0	115.0
Washington	56.0	34.0	97.0	265.0
West Virginia	0.0	0.0	0.0	3.0
Wisconsin	12.0	3.0	10.0	66.0
Wyoming	0.0	0.0	0.0	1.0

Melt

```
melted_result = pd.melt(pivot_table_result.reset_index(), id_vars=["State"], var_name="Ship Mode", value_name="Order Count")
print(melted_result)
```

	State	Ship Mode	Order Count
0	Alabama	First Class	9.0
1	Arizona	First Class	42.0
2	Arkansas	First Class	10.0
3	California	First Class	302.0
4	Colorado	First Class	43.0
..
191	Virginia	Standard Class	115.0
192	Washington	Standard Class	265.0
193	West Virginia	Standard Class	3.0
194	Wisconsin	Standard Class	66.0
195	Wyoming	Standard Class	1.0

[196 rows x 3 columns]

[Q9] What is the advantage of using `melt`?

Ans: it helps transform "wide" format data into "long" format data, making it more suitable for certain types of analysis and visualization

[Q10] From the `superstore_order`, display the ascending order considering values in the 'Profit' column to group the 'Category'.

```
#enter your code
# Group by Category, sum the Profit, and sort in ascending order
```



```
result = superstore_order.groupby('Category')['Profit'].sum().sort_values(ascending=True)
print(result)
```

```
Category
Furniture      16858.5619
Office Supplies 105827.0238
Technology     133410.4932
Name: Profit, dtype: float64
```

[Q11] Create a new column that calculates the total price (sale*quantity) before discount then group by 'product id' and 'category', then show the mean of the total price

```
#enter your code here
# Create a new column for total price before discount
superstore_order['Total_Price'] = superstore_order['Sales'] * superstore_order['Quantity']

# Group by product id and category, then calculate mean of total price
result = superstore_order.groupby(['Product ID', 'Category'])['Total_Price'].mean()
print(result)
```

```
Product ID      Category
FUR-BO-10000112  Furniture      7426.566000
FUR-BO-10000330  Furniture     1258.192000
FUR-BO-10000362  Furniture     1726.898000
FUR-BO-10000468  Furniture      426.532400
FUR-BO-10000711  Furniture     3194.100000
...
TEC-PH-10004912  Technology     747.320000
TEC-PH-10004922  Technology     673.249500
TEC-PH-10004924  Technology      57.149333
TEC-PH-10004959  Technology     412.009000
TEC-PH-10004977  Technology     2441.475429
Name: Total_Price, Length: 1846, dtype: float64
```

[Q12] Complete the function to apply ratio column that calculates from First Class and Standard Class columns on pivot_table_result

```
# [Q12] Complete the function to apply `ratio` column that calculates from `First Class` and `Standard Class` columns on `pivot_table_result`

# function to transform the ratio
def get_class_ratio(row):

    # get the first class column
    first_class = row['First Class']

    # get the standard class column
    standard_class = row['Standard Class']

    # calculate the ratio
    ratio = first_class / standard_class

    return ratio

pivot_table_result["ratio"] = pivot_table_result.apply(get_class_ratio, axis=1)

pivot_table_result.head()
```

```
Ship Mode  First Class  Same Day  Second Class  Standard Class  ratio
State
Alabama      9.0        1.0        18.0          30.0  0.300000
Arizona     42.0        15.0        22.0         123.0  0.341463
Arkansas     10.0         2.0         8.0          35.0  0.285714
California   302.0       106.0       346.0        1000.0  0.302000
Colorado     43.0         5.0        32.0          95.0  0.452632
```

[Q13] After complete Q12, What does the apply function do?

Ans: to perform operations that are more complex than simple vectorized operations and returns a new Series with the results of applying the function to each row/column


[Q14] Create a new column(short_ratio) that works the same as Q12 but with lambda function

```

pivot_table_result["short_ratio"] = pivot_table_result.apply(lambda row: row['First Class'] / row['Standard Class'], axis=1)

pivot_table_result.head()

```



Ship Mode	First Class	Same Day	Second Class	Standard Class	ratio	short_ratio
State						
Alabama	9.0	1.0	18.0	30.0	0.300000	0.300000
Arizona	42.0	15.0	22.0	123.0	0.341463	0.341463
Arkansas	10.0	2.0	8.0	35.0	0.285714	0.285714
California	302.0	106.0	346.0	1000.0	0.302000	0.302000
Colorado	43.0	5.0	32.0	95.0	0.452632	0.452632

[Q15] What is the difference between using `function` in `apply` and `lambda` function? give 2 examples use case.

Ans: Key differences: Regular functions are better for complex logic, multiple operations, or reusable code and documented and are more readable for complex operations Lambda functions are better for one-line operations, anonymous and can't be reused elsewhere in the code

Use cases: Use regular functions when I need documentation, complex logic, or reusability Use lambda functions when the operation is simple and used only once